February, 11, 2026

## 📌 Part 1 – Question Bank

### 10-Mark Questions

1. **Explain the architecture and core benefits of the Logical Volume Manager (LVM) in Linux.**
2. **Describe the step-by-step process of creating, formatting, and persistently mounting a disk partition using command-line tools.**
3. **Provide an in-depth analysis of major Linux filesystems and their specific use cases.**
4. **Compare and contrast the MBR and GPT partitioning schemes in detail.**

### 5-Mark Questions

5. **Explain the procedures for extending and reducing a Logical Volume (LV) in LVM.**
6. **Discuss the Linux Filesystem Hierarchy Standard (FHS) and its essential directories.**
7. **What is Swap Space, and how can a system administrator manage a swap file?**
8. **Compare the primary command-line partitioning tools: fdisk, gdisk, and parted.**
9. **Explain the concept of Journaling and its importance in modern filesystems.**
10. **Outline the best practices for partitioning a Linux system for optimal performance and recovery.**

### 2-Mark Questions

11. **Define Physical Volume (PV) in the context of LVM.**
12. **What is the purpose of the /etc/fstab file?**
13. **Briefly explain the utility of the blkid command.**
14. **Define a Logical Volume (LV).**
15. **What is the function of the TRIM command for SSDs?**

## EXAM POV NOTES (Q&A FORMAT)

📌 **Part 2 – Complete Answer Bank**
**10-Mark Questions**

**Q1. Explain the architecture and core benefits of the Logical Volume Manager (LVM) in Linux.**
**Introduction: LVM** is a sophisticated storage management system that provides a layer of abstraction between physical storage and the operating system, allowing for flexible disk management.

- **Dynamic Resizing Ability:** LVM allows administrators to **resize logical volumes (LVs)** on the fly without the need to unmount the filesystem in many cases. This flexibility ensures that storage can grow or shrink based on real-time data needs without causing system downtime.
- **Storage Pooling Concept:** LVM combines multiple physical disks or partitions into a single large storage pool known as a **Volume Group (VG)**. This allows for the efficient aggregation of diverse storage resources into a unified management unit.
- **Snapshot Functionality:** The system can create read-only or writable **snapshots** of a volume at a specific point in time. These snapshots are invaluable for performing consistent backups or testing changes without affecting the original production data.
- **Efficient Space Utilization:** Unlike rigid traditional partitions, LVM allocates space to volumes only as needed. This prevents the common issue of having one partition completely full while another has significant wasted, unused space.
- **Performance via Striping:** LVM supports data **striping**, which spreads data across multiple physical disks to improve I/O performance. By reading and writing to several disks simultaneously, the system can achieve much higher throughput than a single disk.
- **Physical Volume (PV) Layer:** A **PV** is the base layer of LVM, consisting of a physical disk or partition initialized for LVM use. The pvcreate command is used to prepare these devices so they can be recognized and managed by the LVM system.
- **Volume Group (VG) Layer:** The **VG** acts as a storage container that pools the capacity of one or more Physical Volumes together. Administrators use the vgcreate command to define these groups, which then serve as the source of space for logical volumes.
- **Logical Volume (LV) Layer: LVs** are virtual partitions created from the free space available in a Volume Group. They are the components that are actually formatted with a filesystem and mounted for user data storage, created via the lvcreate command.
- **Improved Reliability through Mirroring:** LVM provides **mirroring** capabilities that keep identical copies of data on different physical disks. This redundancy protects the system from data loss in the event of a single physical disk failure.
- **Scalability for Large Systems:** Because LVs can span across multiple physical disks, they can be much larger than any single physical drive. This makes LVM essential for enterprise environments managing massive datasets across many storage arrays.

**Q2. Describe the step-by-step process of creating, formatting, and persistently mounting a disk partition using command-line tools. Introduction:** Disk partitioning involves dividing a storage device into separate sections, which is a required step before a disk can be used for data storage in Linux.

- **Identify Available Disks:** The first step involves listing all available storage devices using commands like lsblk or sudo fdisk -l. This allows the administrator to identify the correct disk identifier, such as /dev/sdb, before making changes.
- **Select Target Device:** Use the fdisk utility followed by the disk path (e.g., sudo fdisk /dev/sdX) to enter the interactive partitioning menu. This ensures that all subsequent commands are applied only to the intended storage device.

- **Initiate New Partition:** Inside the interactive menu, typing the n command signals the creation of a **new partition**. The system will then prompt the user to define the partition as either a primary or an extended type.
- **Define Partition Parameters:** The administrator must specify a partition number and the physical size of the volume. Size can often be defined in sectors or megabytes to ensure the disk space is allocated precisely according to system requirements.
- **Commit Changes to Disk:** Partitioning changes are only stored in memory until the w command is issued to **write** the table to the disk. This finalizes the layout and exits the utility, making the new partition visible to the kernel.
- **Format with a Filesystem:** A raw partition must be formatted with a filesystem like **ext4** using the sudo mkfs.ext4 /dev/sdXn command. This creates the necessary data structures for storing and retrieving files on that specific section of the disk.
- **Create a Mount Point:** Use sudo mkdir /mnt/new_partition to create a directory that will serve as the entry point for the new storage. This directory acts as the bridge between the Linux directory tree and the physical disk contents.
- **Manual Mount for Testing:** The sudo mount /dev/sdXn /mnt/new_partition command is used to manually attach the filesystem to the directory tree. This allows the administrator to verify that the partition is accessible and the formatting was successful.
- **Retrieve Unique Identifier (UUID):** For persistent mounting, it is best practice to find the **UUID** using the sudo blkid command. Using the UUID instead of the device path prevents mounting errors if disk letters change after a hardware update.
- **Configure Persistent Mounting:** Finally, the UUID and mount parameters are added to the **/etc/fstab** configuration file. This ensures the operating system automatically mounts the partition every time the system boots up.

**Q3. Provide an in-depth analysis of major Linux filesystems and their specific use cases. Introduction:** A **filesystem** is a method for organizing and managing files on a storage device, with Linux supporting various types optimized for different workloads.

- **The ext4 Filesystem:** As the modern standard for Linux, **ext4** is a journaling filesystem that supports very large volumes and efficient data management. It is the recommended choice for general Linux desktop and server usage.
- **XFS for High Performance: XFS** is a 64-bit filesystem designed for extreme scalability and high-performance parallel I/O operations. It is best suited for large-scale data storage environments where speed and reliability are critical.
- **Btrfs (B-tree Filesystem):** This modern filesystem offers advanced features like **snapshotting**, data deduplication, and integrated device management. It is often used in modern distributions for its flexibility in managing storage pools.
- **ZFS for Enterprise Reliability: ZFS** provides enterprise-grade features including built-in RAID support, high data integrity, and checksumming. It is the primary choice for servers where preventing data corruption is the highest priority.
- **F2FS for Flash Storage:** The **Flash-Friendly File System (F2FS)** is specifically optimized for the characteristics of SSDs and mobile storage. It improves performance and longevity for devices that rely on NAND flash memory.
- **NTFS for Windows Compatibility:** While native to Windows, Linux supports **NTFS** for external drives or dual-boot partitions. It allows for metadata journaling and large file support while maintaining cross-platform access.

- **FAT32 and exFAT Utility:** These legacy filesystems are primarily used for USB drives and SD cards due to their broad compatibility across all operating systems. However, they lack modern features like journaling and are prone to corruption.
- **The Role of Journaling:** Filesystems like ext4 and XFS use **journaling** to keep a log of changes, which prevents corruption during power failures. This feature is a key differentiator between modern robust filesystems and legacy ones.
- **Scalability and Addressing:** Modern 64-bit filesystems like XFS and Btrfs allow for addressing massive amounts of storage that older 32-bit systems cannot handle. This makes them essential for modern servers with multi-terabyte drives.
- **Selection Criteria:** Choosing the right filesystem depends on the specific balance of **security, performance, and compatibility** required for the task. Factors like robustness and accessibility play a vital role in this selection process.

**Q4. Compare and contrast the MBR and GPT partitioning schemes in detail. Introduction:**
Partitioning schemes define how information about partitions is stored on a disk, determining limits on size and the number of partitions.

| Feature | MBR (Master Boot Record) | GPT (GUID Partition Table) |
|---|---|---|
| **Partition Limit** | Supports only up to 4 primary partitions. | Supports a practically unlimited number of partitions. |
| **Max Disk Size** | Limited to a maximum of 2TB disk size. | Supports disks significantly larger than 2TB. |
| **Addressing** | Uses a 32-bit partition table to save addresses. | Uses a 64-bit partition table for expanded addressing. |
| **Identification** | Identifies partitions using simple numerical values. | Uses Universally Unique Identifiers (GUIDs). |
| **Firmware Standard** | Designed for legacy PC BIOS systems. | Part of the modern UEFI standard replacement for BIOS. |

- **MBR Extended Partitions:** To bypass the four-partition limit, MBR uses **extended partitions** which act as containers for multiple logical partitions. This adds complexity to the partition layout compared to the flat structure of GPT.
- **GPT Data Redundancy:** GPT stores multiple copies of the partition table at different locations on the disk. This provides built-in **redundancy**, allowing the system to recover the layout if the primary table is corrupted.
- **GPT Modern Recommendation:** Because of its support for large disks and data integrity features, GPT is the **recommended standard** for all modern systems. It is necessary for booting most modern 64-bit operating systems.
- **MBR Legacy Support:** MBR remains relevant for older 32-bit systems and legacy BIOS hardware that cannot recognize GPT. It is widely compatible but technically obsolete for modern high-capacity drives.
- **Disk Structure Representation:** Linux represents these partitions hierarchically in the /dev directory, such as /dev/sda1 for MBR or GPT volumes. Regardless of the scheme, the OS treats them as independent storage units once initialized.

**5-Mark Questions**

**Q5. Explain the procedures for extending and reducing a Logical Volume (LV) in LVM.**
**Introduction:** LVM provides the ability to manage volume sizes dynamically to meet changing storage requirements.

- **Extending the LV:** The lvextend command is used to add space to an existing volume, such as lvextend -L +5G /dev/my_vg/my_lv. This increases the logical capacity of the volume within the volume group.
- **Resizing the Filesystem:** After extending the LV, the underlying filesystem must be expanded using resize2fs to recognize the new space. This step ensures the extra capacity is actually usable by the operating system.
- **Reducing the LV - Unmounting:** Unlike extension, reducing a volume requires the partition to be **unmounted** first using umount /mnt. This prevents data corruption while the volume size is being modified.
- **Executing the Reduction:** The lvreduce command is then used to shrink the volume to the desired size, such as lvreduce -L 5G /dev/my_vg/my_lv. Care must be taken to ensure the new size is not smaller than the data currently stored.
- **Components Removal:** If a volume is no longer needed, it can be removed entirely using lvremove, followed by removing the VG and PV if necessary. This process completely deallocates the space back to the physical disk.

**Q6. Discuss the Linux Filesystem Hierarchy Standard (FHS) and its essential directories.**
**Introduction:** The **FHS** defines a tree-like structure for organizing files and directories in Linux to ensure consistency across distributions.

- **The /bin Directory:** This folder contains **essential system binaries** and executable programs required for the system to function. It includes basic commands like ls, cp, and mv.
- **The /boot Directory:** This directory holds all **bootloader files**, including the Linux kernel and configuration needed to start the computer. It is critical for the initial startup phase of the OS.
- **The /etc Directory:** This is the central location for all **system-wide configuration files**. It contains scripts and settings that control how services and applications behave on the system.
- **The /home Directory:** This section is dedicated to **user directories**, where personal files and settings are stored. Keeping this separate from the system root is a best practice for easier recovery.
- **The /var and /tmp Directories:** /var is used for **variable data** like logs and caches, while /tmp holds temporary files created by applications. These directories handle frequently changing data that does not need to be persistent.

**Q7. What is Swap Space, and how can a system administrator manage a swap file?**
**Introduction: Swap space** is used as virtual memory when the physical RAM is full, allowing inactive memory pages to be moved to the disk.

- **Create the Swap File:** An administrator can create a swap file of a specific size using the fallocate command. For example, sudo fallocate -l 2G /swapfile reserves 2GB of disk space for virtual memory.
- **Secure File Permissions:** It is essential to restrict access to the swap file using sudo chmod 600 /swapfile. This ensures that only the root user can read the sensitive memory data stored within the file.

- **Initialize Swap Area:** The file must be formatted as a swap area using the **mkswap** command. This prepares the file structure so the Linux kernel can recognize it as a valid swap space.
- **Activate Swap File:** The swapon command is used to enable the file for immediate use by the system. Usage can then be monitored using commands like free -h or swapon --show.
- **Make Swap Permanent:** To ensure the swap is available after rebooting, an entry must be added to **/etc/fstab**. This automates the activation of the swap file during the boot process.

**Q8. Compare the primary command-line partitioning tools: fdisk, gdisk, and parted.**
**Introduction:** Linux provides several command-line utilities for managing disk partitions, each with specific strengths and compatibility.

- **fdisk Utility:** This is a classic CLI tool primarily used for managing **MBR partition tables**. While it is very common, it is generally not used for modern GPT disks or drives larger than 2TB.
- **gdisk Utility:** Specifically designed as the **GPT counterpart to fdisk**, gdisk is used for managing GUID Partition Tables. it offers powerful features for modern partitioning needs on newer hardware.
- **parted Utility:** This is a highly flexible tool that supports **both MBR and GPT** partitioning schemes. It is often used for script-based partitioning and supports resizing and moving partitions directly.
- **Interactive vs. Command-based:** While fdisk and gdisk use an interactive menu system, parted can be used both interactively and with direct command-line arguments. This makes parted more suitable for automation.
- **Selection Logic:** Administrators choose fdisk for legacy systems, gdisk for modern GPT disks, and parted when they need a single tool that handles all formats or complex resizing.

**Q9. Explain the concept of Journaling and its importance in modern filesystems.**
**Introduction: Journaling** is a feature in modern filesystems like ext4 and XFS that helps prevent data corruption by keeping a log of changes.

- **The Journal Log:** The filesystem maintains a dedicated log (journal) where it records intended file changes before they are permanently written to the disk. This acts as a "to-do list" for the filesystem.
- **Crash Recovery:** In the event of a power failure or system crash, the OS can check the journal to see which operations were incomplete. It can then either finish those tasks or discard them to maintain consistency.
- **Integrity Maintenance:** Journaling significantly reduces the risk of **data loss** compared to older filesystems like FAT32. It ensures that the metadata of the filesystem remains in a valid state.
- **Speed of Repair:** Without a journal, the system must perform a time-consuming full disk check (fsck) after a crash. A journaling filesystem only needs to replay the log, which is much faster.
- **Standard Adoption:** Most modern Linux filesystems, including **ext3, ext4, XFS, and Btrfs**, utilize journaling by default. This has made Linux systems much more robust for mission-critical server environments.

**Q10. Outline the best practices for partitioning a Linux system for optimal performance and recovery.**

- **Introduction:** Proper disk partitioning is essential for organizing data, improving performance, and implementing security protocols.

- **Use GPT for New Systems:** It is recommended to use **GPT** for all new installations and large disks. GPT provides better reliability through redundancy and supports modern hardware standards.
- **Separate Root and Home:** Keeping the **/ (root)** and **/home** directories on separate partitions is a major best practice. This allows for easier system reinstallation or recovery without losing personal user data.
- **Allocate Sufficient Swap:** Ensure that swap space is allocated, typically sized between **1x to 2x the amount of RAM**. This provides a safety net for the system during periods of high memory usage.
- **Avoid Excessive Partitions:** While organization is good, creating too many small partitions can lead to wasted space and management complexity. Partition only as much as necessary for data segregation.
- **SSD Optimization:** For systems using SSDs, it is important to ensure **partition alignment** and enable the **TRIM** command. These steps optimize the performance and longevity of flash-based storage devices.

**2-Mark Questions**

**Q11. Define Physical Volume (PV) in the context of LVM.**

- **Definition:** A **PV** is a physical disk or partition that has been initialized for use with LVM using the pvcreate command.
- **Function:** It serves as the fundamental building block that provides the raw storage capacity to a Volume Group.

**Q12. What is the purpose of the /etc/fstab file?**

1. **Definition:** The **/etc/fstab** file is a system configuration file used to define how disk partitions and filesystems are mounted.
2. **Purpose:** It ensures that specified storage devices are automatically and persistently mounted at the correct locations during every system boot.

**Q13. Briefly explain the utility of the blkid command.**

- **Definition:** The **blkid** command is a command-line utility used to locate and display the attributes of block devices.
- **Utility:** It is primarily used by administrators to find the **UUID** of a partition, which is required for secure and persistent mounting in the fstab file.

**Q14. Define a Logical Volume (LV).**

- **Definition:** An **LV** is a virtual partition created from the pooled space in a Volume Group.
- **Usage:** It acts like a traditional partition that can be formatted with a filesystem and mounted for actual data storage.

**Q15. What is the function of the TRIM command for SSDs?**

- **Definition: TRIM** is a command (managed in Linux via fstrim) that informs an SSD which blocks of data are no longer in use.
- **Function:** It helps maintain the **performance and longevity** of the SSD by allowing the drive to handle garbage collection more efficiently.